

Amazon Fine Food Reviews using t-SNE

August 13, 2018

1 t-SNE Visualization of Amazon reviews with polarity based color-coding

BY: Subham Sarkar

```
In [3]: %matplotlib inline

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re

import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
```

```

# using the SQLite Table to read data.
con = sqlite3.connect('C:/Users/Subham Sarkar/Downloads/database.sqlite')

#filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)

# Give reviews with Score>3 a positive rating, and reviews with a score<3 a negative rating
def partition(x):
    if x < 3:
        return 'negative'
    return 'positive'

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative

import warnings
warnings.filterwarnings("ignore")

```

```

In [4]: print(filtered_data.shape) #looking at the number of attributes and size of the data
filtered_data.head()

```

```

(525814, 10)

```

```

Out[4]:

```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	d11 pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	positive	1303862400	
1	0	0	negative	1346976000	
2	1	1	positive	1219017600	
3	3	3	negative	1307923200	
4	0	0	positive	1350777600	

Summary

Text

```

0 Good Quality Dog Food I have bought several of the Vitality canned d...
1 Not as Advertised Product arrived labeled as Jumbo Salted Peanut...
2 "Delight" says it all This is a confection that has been around a fe...
3 Cough Medicine If you are looking for the secret ingredient i...
4 Great taffy Great taffy at a great price. There was a wid...

```

```

In [5]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()

```

```

Out[5]:      Id  ProductId      UserId      ProfileName  HelpfulnessNumerator \
0   78445  B000HDL1RQ  AR5J8UI46CURR  Geetha Krishnan             2
1  138317  B000HDOPYC  AR5J8UI46CURR  Geetha Krishnan             2
2  138277  B000HDOPYM  AR5J8UI46CURR  Geetha Krishnan             2
3   73791  B000HDOPZG  AR5J8UI46CURR  Geetha Krishnan             2
4  155049  B000PAQ75C  AR5J8UI46CURR  Geetha Krishnan             2

```

```

      HelpfulnessDenominator  Score      Time \
0                2          5  1199577600
1                2          5  1199577600
2                2          5  1199577600
3                2          5  1199577600
4                2          5  1199577600

```

```

      Summary \
0  LOACKER QUADRATINI VANILLA WAFERS
1  LOACKER QUADRATINI VANILLA WAFERS
2  LOACKER QUADRATINI VANILLA WAFERS
3  LOACKER QUADRATINI VANILLA WAFERS
4  LOACKER QUADRATINI VANILLA WAFERS

```

```

      Text
0  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

```

In [6]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

```

```

In [7]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep=False)
final.shape

```

```
Out[7]: (364173, 10)
```

```
In [8]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

```
Out[8]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2VOI904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0		3	1	5	1224892800
1		3	2	4	1212883200

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [9]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [10]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(364171, 10)
```

```
Out[10]: positive    307061
negative      57110
Name: Score, dtype: int64
```

```
In [11]: # find sentences containing HTML tags
import re
i=0;
for sent in final['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;
```

6

I set aside at least an hour each day to read to my son (3 y/o). At this point, I consider myself

```
In [12]: nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to C:\Users\Subham
```

```
[nltk_data]      Sarkar\AppData\Roaming\nltk_data...
```

```
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[12]: True
```

```
In [13]: stop = set(stopwords.words('english')) #set of stopwords
```

```
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
```

```
def cleanhtml(sentence): #function to clean the word of any html-tags
```

```
    cleanr = re.compile('<.*?>')
```

```
    cleantext = re.sub(cleanr, ' ', sentence)
```

```
    return cleantext
```

```
def cleanpunc(sentence): #function to clean the word of any punctuation or special char
```

```
    cleaned = re.sub(r'[?|!|\\'|"|#]', '', sentence)
```

```
    cleaned = re.sub(r'[.,|)|(|\\|/]', '', cleaned)
```

```
    return cleaned
```

```
print(stop)
```

```
print('*****')
```

```
print(sno.stem('tasty'))
```

```
{'nor', 'between', 'am', 'has', 'out', 'by', 'too', 'theirs', 'once', 'only', 'was', 'very', 'tast
```

```
*****
```

```
tasti
```

```
In [14]: #Code for implementing step-by-step the checks mentioned in the pre-processing phase
```

```
# this code takes a while to run as it needs to run on 500k sentences.
```

```
i=0;
```

```
str1=' ';
```

```
final_string=[];
```

```
all_positive_words=[] # store words from +ve reviews here
```

```
all_negative_words=[] # store words from -ve reviews here.
```

```
s=''
```

```
for sent in final['Text'].values:
```

```
    filtered_sentence=[]
```

```
    #print(sent);
```

```
    sent=cleanhtml(sent) # remove HTML tags
```

```
    for w in sent.split():
```

```
        for cleaned_words in cleanpunc(w).split():
```

```
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
```

```
                if(cleaned_words.lower() not in stop):
```

```

        s=(sno.stem(cleaned_words.lower())).encode('utf8')
        filtered_sentence.append(s)
        if (final['Score'].values)[i] == 'positive':
            all_positive_words.append(s) #list of all words used to descr
        if (final['Score'].values)[i] == 'negative':
            all_negative_words.append(s) #list of all words used to descr
        else:
            continue
    else:
        continue
    #print(filtered_sentence)
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    #print("*****")

    final_string.append(str1)
    i+=1

```

In [15]: `print(final_string)`

IOPub data rate exceeded.

The notebook server will temporarily stop sending output

to the client in order to avoid crashing it.

To change this limit, set the config variable

`--NotebookApp.iopub_data_rate_limit`.`

Current values:

`NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)`

`NotebookApp.rate_limit_window=3.0 (secs)`

In [16]: `final['CleanedText']=final_string #adding a column of CleanedText which displays the`
`final['CleanedText']=final['CleanedText'].str.decode("utf-8")`

In [17]: `final.head(3) #below the processed review can be seen in the CleanedText Column`

```

# store final table into an SQLite table for future.
conn = sqlite3.connect('final.sqlite')
c=conn.cursor()
conn.text_factory = str
final.to_sql('Reviews', conn, schema=None, if_exists='replace', index=True, index_label=

```

2 Bag of Words

In [18]: `#BoW`

```

count_vect = CountVectorizer() #in scikit-learn
final_counts = count_vect.fit_transform(final['CleanedText'].values)

```

```

print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (364171, 71624)
the number of unique words 71624

```

```

In [19]: freq_dist_positive=nlTK.FreqDist(all_positive_words)
         freq_dist_negative=nlTK.FreqDist(all_negative_words)
         print("Most Common Positive Words : ",freq_dist_positive.mOST_common(20))
         print("Most Common Negative Words : ",freq_dist_negative.mOST_common(20))

```

```

Most Common Positive Words : [(b'like', 139429), (b'tast', 129047), (b'good', 112766), (b'fla
Most Common Negative Words : [(b'tast', 34585), (b'like', 32330), (b'product', 28218), (b'one

```

3 Observation:-

From the above it can be seen that the most common positive and the negative words overlap for eg. 'like' could be used as 'not like' etc. So, it is a good idea to consider pairs of consequent words (bi-grams) or q sequence of n consecutive words (n-grams)

```

In [20]: import warnings
         warnings.filterwarnings('ignore')
         from sklearn.preprocessing import StandardScaler

         standardized_data=StandardScaler().fit_transform(final_counts[4000].toarray())
         print(standardized_data.shape)

```

```

(1, 71624)

```


4 t-SNE plot for Bag of Words with perplexity = 30

```

In [30]: from sklearn.manifold import TSNE
         label_4000 = final['Score'][0:4000]

         model = TSNE(n_components=2, random_state=0)
         tsne_data = model.fit_transform(standardized_data)
         tsne_data=np.vstack((tsne_data.T,label_4000)).T
         tsne_df=pd.DataFrame(data=tsne_data,columns=('Dim_1','Dim_2','Score'))
         sns.FacetGrid(tsne_df, hue='Score', size=6).map(plt.scatter, 'Dim_1', 'Dim_2').add_le
         plt.show()

```



5 tf-idf (term frequency-inverse document frequency)

```
In [21]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2))
         final_tf_idf = tf_idf_vect.fit_transform(final['CleanedText'].values)
         print("the type of count vectorizer ",type(final_tf_idf))
         print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
         print("the number of unique words including both unigrams and bigrams ", final_tf_idf
```

```
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer  (364171, 2923725)
the number of unique words including both unigrams and bigrams  2923725
```

```
In [22]: features = tf_idf_vect.get_feature_names()
         print("some sample features(unique words in the corpus)",features[100000:100010])

some sample features(unique words in the corpus) ['antler fair', 'antler fall', 'antler fantas
```



```
In [28]: print(final_tf_idf[0:100,:].toarray()[0])
```

```
[0. 0. 0. ... 0. 0. 0.]
```

```
In [29]: def top_tfidf_feats(row, features, top_n=25):
```

```
    ''' Get top n tfidf values in row and return them with their corresponding features '''
    topn_ids = np.argsort(row)[::-1][:top_n]
    top_feats = [(features[i], row[i]) for i in topn_ids]
    df = pd.DataFrame(top_feats)
    df.columns = ['feature', 'tfidf']
    return df
```

```
top_tfidf = top_tfidf_feats(final_tf_idf[1,:].toarray()[0], features, 25)
```

```
In [30]: top_tfidf
```

```
Out[30]:
```

	feature	tfidf
0	page open	0.192673
1	read sendak	0.192673
2	movi incorpor	0.192673
3	paperback seem	0.192673
4	version paperback	0.192673
5	flimsi take	0.192673
6	incorpor love	0.192673
7	rosi movi	0.192673
8	keep page	0.192673
9	grew read	0.192673
10	realli rosi	0.186715
11	sendak book	0.186715
12	cover version	0.186715
13	miss hard	0.182488
14	kind flimsi	0.176530
15	hard cover	0.174265
16	watch realli	0.170572
17	book watch	0.170572
18	sendak	0.166345
19	howev miss	0.165169
20	paperback	0.162117
21	hand keep	0.153894
22	two hand	0.150202
23	rosi	0.148291
24	seem kind	0.147253

```
In [31]: import warnings
```

```
warnings.filterwarnings('ignore')
```

```
from sklearn.preprocessing import StandardScaler
```

```
standardize_data=StandardScaler().fit_transform(final_tf_idf[0:100,:].toarray())
```

```
print(standardize_data.shape)
```

$$(100, 2923725)$$

6 t-SNE plot of average tf-idf with perplexity = 30

```
In [ ]: from sklearn.manifold import TSNE
        label_4000 = final['Score'][0:4000]

        model = TSNE(n_components=2, random_state=0)
        tsne_data = model.fit_transform(standardize_data)
        tsne_data=np.vstack((tsne_data.T,label_4000)).T
        tsne_df=pd.DataFrame(data=tsne_data,columns=('Dim_1','Dim_2','Score'))
        sns.FacetGrid(tsne_df, hue='Score', size=6).map(plt.scatter, 'Dim_1', 'Dim_2').add_legend()
        plt.show()
```

7 Word2Vec

```
In [33]: i=0
         list_of_sent=[]
         for sent in final['CleanedText'].values:
             list_of_sent.append(sent.split())

In [34]: print(final['CleanedText'].values[0])
         print("*****")
         print(list_of_sent[0])
```

```
witti littl book make son laugh loud recit car drive along alway sing refrain hes learn whale :
*****
['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'car', 'drive', 'along', 'a
```

```
In [35]: w2v_model=Word2Vec(list_of_sent,min_count=5,size=50, workers=4)
```

```
In [36]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times  21938
sample words  ['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'car', 'driv
```

```
In [37]: w2v_model.wv.most_similar('tasti')
```

```
Out[37]: [('delici', 0.8079195618629456),
          ('yummi', 0.7775864601135254),
          ('tastey', 0.722675621509552),
          ('good', 0.6968026161193848),
          ('nice', 0.6770238876342773),
```

```
(('hearti', 0.6712539196014404),
 ('satisfi', 0.6701481938362122),
 ('nutriti', 0.644965648651123),
 ('great', 0.6448278427124023),
 ('terrifi', 0.6447775363922119])
```

```
In [38]: w2v_model.wv.most_similar('like')
```

```
Out[38]: [('weird', 0.7477270364761353),
 ('dislik', 0.7056392431259155),
 ('funke', 0.6866582632064819),
 ('yucki', 0.6820259690284729),
 ('okay', 0.6809946894645691),
 ('prefer', 0.6613717079162598),
 ('gross', 0.6539754271507263),
 ('fake', 0.6539649963378906),
 ('resembl', 0.6470655202865601),
 ('appeal', 0.6464338898658752)]
```

```
In [39]: count_vect_feat = count_vect.get_feature_names() # list of words in the BoW
print(count_vect_feat[count_vect_feat.index('like')])
```

```
like
```

8 TF-IDF weighted Word2Vec

```
In [ ]: # TF-IDF weighted Word2Vec
tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
```

9 average Word2Vec

```
In [49]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))
```

364171

50

10 t-SNE plot of average Word2Vec with perplexity = 30

```
In [54]: from sklearn.manifold import TSNE
label_4000_avg = final['Score'][0:4000]

model = TSNE(n_components=2, random_state=0)
tsne_data = model.fit_transform(sent_vectors[0:4000])
tsne_data=np.vstack((tsne_data.T,label_4000_avg)).T
tsne_df=pd.DataFrame(data=tsne_data,columns=('Dim_1','Dim_2','Score'))
sns.FacetGrid(tsne_df, hue='Score', size=6).map(plt.scatter, 'Dim_1', 'Dim_2').add_le
plt.show()
```

